

Dealing with Separation in Logistic Regression Models^{*}

Carlisle Rainey[†]

Working paper: Comments welcome!

July 21, 2015

ABSTRACT

When facing small numbers of observations or “rare events,” political scientists often encounter separation, in which important explanatory variables perfectly predict binary events or non-events. In this situation, maximum likelihood provides implausible estimates and the researcher must incorporate some form of prior information into the model. The most sophisticated research uses Jeffreys’ invariant prior to stabilize the estimates. While Jeffreys’ prior has the advantage of being automatic, I show that it often provides too much prior information, producing point estimates that are too small and confidence intervals that are too narrow. To help researchers choose a more reasonable prior distribution, I introduce the concept of a partial prior predictive distribution and develop the tools required to compute the partial prior predictive distribution, estimate the subsequent model, and summarize the results.

^{*}I thank Mark Bell and Nicholas Miller for making their data available. I thank Mark Bell, David Firth, Nicholas Miller, and seminar participants at the University of Kansas and Texas A&M University for helpful comments. The analyses presented here were conducted with R 3.2.1 and JAGS 3.4.0. All data and computer code necessary for replication are available at github.com/carlislerainey/priors-for-separation. The (working) R package `separation` implements the procedures discussed in the paper and is available at github.com/carlislerainey/separation.

[†]Carlisle Rainey is Assistant Professor of Political Science, Texas A&M University, 2010 Allen Building, College Station, TX, 77843 (crainey@tamu.edu).

Separation, in which an explanatory variable perfectly predicts some binary observations, remains a problem in political science research (e.g., Bell and Miller 2015; Mares 2015; Vining, Wilhelm, and Collens 2015; Brown and Kaplow 2014; Barrilleaux and Rainey 2014; Leeman and Mares 2014; Reiter 2014; Weisiger 2014; Minozzi and Volden 2013; Cederman, Gleditsch, and Hug 2013; Fuhrmann 2012; Rocca, Sanchez, and Morin 2011; Peterson and Drury 2011; Cox, Kousser, and McCubbins 2010; Ahlquist 2010; Rauchhaus 2009; Smith and Fridkin 2008; Casellas 2009; Desposato and Scheiner 2008; Heller and Mershon 2008; and DeRouen and Bercovitch 2008), Zorn (2005) offers the most principled solution to the problem of separation by suggesting that researchers maximize a penalized version of the usual likelihood function. Zorn's approach has the advantage of being automatic and easy for researchers to use.

However, when implementing Zorn's approach, substantive researchers face two major problems. First, the usual asymptotic confidence intervals and p -values do not work well. While a good method exists for finding confidence intervals and p -values for the model coefficients, this method does not extend to the typical quantities of interest, such as first differences. Researchers must still rely on the (poor) asymptotic approximation to simulate these quantities (King, Tomz, and Wittenberg 2000). Secondly, and perhaps most importantly, while the penalty suggested by Zorn has the attractive property of bias reduction in logistic regression models (Firth 1993), it does not necessarily approximate a researcher's prior information. Some shrinkage toward zero is required to obtain finite estimates, but the appropriate amount of shrinkage depends on the substantive problem and the prior information. To address these two problems, I suggest that researchers use a range of priors, focusing on an informative prior, and use MCMC to simulate directly from the posterior.

In this paper, I introduce conceptual and computational tools that help researchers understand the information provided by a given prior distribution and use that prior distribution to obtain meaningful point estimates and confidence intervals. I make three specific contributions. First, I use statistical theory and two applied examples to demonstrate the importance of choos-

ing a prior distribution that represents actual prior information (and conducting robustness checks using a variety of prior distributions). Second, I introduce the concept of a partial prior predictive distribution, a powerful tool in understanding and choosing a prior when facing separation. Third, I introduce new software that make it easy for researchers to choose an informative prior distribution, simulate directly from the posterior distribution, and summarize the inferences.¹

I begin with a basic overview of the logistic regression model and summary of the impact of separation on the maximum likelihood estimates. I then describe two default prior distributions that researchers might use to handle separation. Next, I use a theoretical result and an applied example to demonstrate the importance of choosing an informative prior. I then introduce researchers to the concept of a partial prior predictive distribution, which enables researchers to understand complex prior distributions in terms of the key quantities of interest. To illustrate how these ideas work in practice, I conclude with a replication of Rauchhaus (2009) and Bell and Miller (2015), whose disagreement about the effect of nuclear weapons on war hinges, in part, on how to deal with separation.

The Logistic Regression Model

Political scientists commonly use logistic regression to model the probability of events such as war (e.g., Fearon 1994), policy adoption (e.g., Berry and Berry 1990), turning out to vote (e.g., Wolfinger and Rosenstone 1980), and government formation (e.g., Martin and Stevenson 2001). In the typical situation, the researcher uses an $n \times (k + 1)$ design matrix X consisting of an intercept and k explanatory variables to model a vector of n binary outcomes y , where $y_i \in \{0, 1\}$, using the model $\Pr(y_i) = \Pr(y_i = 1 \mid X_i) = \frac{1}{1 + e^{-X_i\beta}}$, where β is a parameter vector of length $k + 1$. Using this model, it is straightforward to calculate the likelihood function

¹A working version of the R package `separation`, which implements the ideas in this paper, is available at github.com/carlislerainey/separation.

$$\Pr(y|\beta) = L(\beta|y) = \prod_{i=1}^n \left[\left(\frac{1}{1 + e^{-X_i\beta}} \right)^{y_i} \left(\frac{1}{1 + e^{-X_i\beta}} \right)^{1-y_i} \right].$$

Researchers routinely obtain the maximum likelihood estimate $\hat{\beta}^{mle}$ of the model parameters β by finding the vector β that maximizes L (i.e., maximizing the likelihood of the observed data). While this approach works quite well in most applications, it fails in a situation known as separation (Zorn 2005).

Separation

Separation occurs in models of binary outcome data when one explanatory variable perfectly predicts zeros, ones, or both.² *Complete separation* occurs when the “problematic” explanatory variable s_i (for separating explanatory variable) perfectly predicts both zeros *and* ones. *Quasicomplete separation* occurs when s_i perfectly predicts either zeros *or* ones, but not both (Albert and Anderson 1984 and Zorn 2005). *Overlap*, the ideal case, occurs when there is no such s_i . When the data overlap, the usual maximum likelihood estimates exist and provide reasonable estimates of parameters. However, under complete or quasicomplete separation, finite maximum likelihood estimates do not exist and the usual method of calculating standard errors fails (Albert and Anderson 1984 and Zorn 2005).

For the binary explanatory variable s_i , complete separation occurs when s_i perfectly predicts *both* zeros *and* ones.³ For example, suppose s_i , such that $y_i = 0$ for $s_i = 0$ and $y_i = 1$ for $s_i = 1$. To maximize the likelihood of the observed data, the “S”-shaped logistic regression curve must assign $\Pr(y_i) = \frac{1}{1+e^{-X_i\beta}} = 0$ when $s_i = 0$ and $\Pr(y_i) = \frac{1}{1+e^{-X_i\beta}} = 1$ when $s_i = 1$. Since the logistic regression curve lies strictly between zero and one, this likelihood cannot be

²Separation can also occur when a *combination* of explanatory variables perfectly predicts zeros, ones, or both, see Lesaffre and Albert (1989). See Geyer (2009) for a much more general view of the concept of separation.

³Albert and Anderson (1984, p. 38) write that complete separation occurs when “there is a vector [of coefficients] that correctly allocates all observations to their group.”

achieved, only approached asymptotically as the coefficient β_s for s_i approaches infinity. Thus, the likelihood function under complete separation is monotonic, which implies that a finite maximum likelihood estimate does not exist.⁴

Quasicomplete separation occurs when s_i perfectly predicts *either* zeros *or* ones.⁵ For example, suppose that, when $s_i = 0$, sometimes $y_i = 1$ and other times $y_i = 0$, but $y_i = 1$ *always* when $s_i = 1$. To maximize the likelihood of the observed data, the “S”-shaped logistic regression curve must assign $\Pr(y_i) = \frac{1}{1+e^{-x_i\beta}} \in (0, 1)$ when $s_i = 0$ and $\Pr(y_i) = \frac{1}{1+e^{-x_i\beta}} = 1$ when $s_i = 1$. Again, since the logistic regression curve lies strictly between zero and one, this likelihood cannot be achieved, only approached asymptotically. Thus, the likelihood function under quasicomplete separation also monotonically increases as the coefficient for s_i increases, which again implies that the maximum likelihood estimate does not exist.

For example, Barrilleaux and Rainey (2014) find that no Democratic governors opposed the Medicaid expansion under the Affordable Care Act, leading to a maximum likelihood estimate of positive infinity for the coefficient for the indicator of Republican governors. Similarly, Rauchhaus (2009) (see Bell and Miller 2015) finds no instances of states with nuclear weapons engaging in war with each other, leading to an estimated coefficient of negative infinity for the coefficient for the variable indicating nuclear dyads. To maximize the likelihood in these situations, the model must assign zero probability of opposition to states with Democratic governors and zero probability of war to nuclear dyads. Because the logistic regression curve lies strictly above zero, this cannot happen, though it can be approached asymptotically as the

⁴Although coefficient estimates under separation are infinite *in theory*, the hill-climbing algorithms approximate the infinite estimates with large, finite values *in practice*. These approximations increase with the precision of the algorithm. See Zorn (2005) for an illustration using software in Stata and R.

⁵I describe complete and quasicomplete separation for a *binary* explanatory variable, which is less general, but more explicable than the general case considered by Albert and Anderson (1984). My approach also follows the convention of Heinze and Schemper (2002) and Zorn (2005). Indeed, in social science problem, binary explanatory variables more commonly lead to separation, so little is lost. But according to Albert and Anderson’s general definition, quasicomplete separation occurs when all cases are perfectly predicted except at a single value x^* of an explanatory (separating) variable x_i . For the binary case, x^* must be either zero or one. Thus, for a binary x_i , quasicomplete separation occurs when x_i perfectly predicts *either* zeros *or* ones.

coefficient for s_i goes to negative infinity.

For simplicity, this paper focuses on quasicomplete separation with binary explanatory variables, the most common situation for social scientists. However, the ideas apply equally well to the less common situation of complete separation and/or continuous explanatory variables. For convenience, I say that the “direction of the separation” is positive if and only if $s_i = 1 \implies y_i = 1$ or $s_i = 0 \implies y_i = 0$ and that the direction of separation is negative if and only if $s_i = 0 \implies y_i = 1$ or $s_i = 1 \implies y_i = 0$. Thus, $\hat{\beta}^{mle} = +\infty$ when the direction of the estimate is positive, and $\hat{\beta}^{mle} = -\infty$ when the direction of the estimate is negative.

Solutions to Separation

The maximum likelihood (ML) framework requires the researcher to find the parameter vector that “maximizes the likelihood of the observed data.” Of course, infinite coefficients *always* generate separated data, while finite coefficients only *sometimes* generate separation. Thus, under separation, the ML can only produce infinite estimates.

Before addressing potential solutions to this problem, let me mention two unsatisfactory “solutions” found in applied work. In some cases, researchers simply ignore the problem of separation and interpret the large estimates and standard errors as though these are reasonable. However, this approach leads researchers to overstate the magnitude of the effect and the uncertainty of the estimates.⁶ Secondly, researchers sometimes “solve” the problem of separation by dropping the separating variable from the model. Zorn (2005, pp. 161-162) correctly

⁶This paper focuses on the situation in which the researcher considers infinite logistic regression coefficients “implausibly large.” Almost all social science work falls into this category. Indeed, for most social science problems, it makes little sense to think that events *can never* happen or *always* happen in certain contexts. There must always be some chance of an event or no event, though perhaps fleetingly small. As Heinze and Schemper (2002, p. 2409) note: “In general, one does not assume infinite parameter values in underlying populations.” For most social science research, then, it makes sense to consider infinite logistic regression coefficient estimates as a quirk of maximum likelihood estimation and not as a reasonable guess of the actual parameter. However, if a researcher wishes to adopt the perspective that infinite coefficient estimates are indeed reasonable, then Geyer (2009) offers a useful perspective.

dismisses this approach:

As a practical matter, separation forces the analyst to choose from a number of problematic alternatives for dealing with the problem. The most widely used “solution” is simply to omit the offending variable or variables from the analysis. In political science, this is the approach taken in a number of studies in international relations, comparative politics, and American politics. It is also the dominant approach in sociology, economics, and the other social sciences, and it is the recommended method in a few prominent texts in statistics and econometrics. Of course, this alternative is a particularly unattractive one; omitting a covariate that clearly bears a strong relationship to the phenomenon of interest is nothing more than deliberate specification bias.

One principled solution is to build prior information $p(\beta)$ (the same prior information that leads researchers to deem infinite coefficients “implausibly large”) into the model using Bayes’ rule, so that

$$p(\beta|y) = \frac{\overbrace{p(y|\beta)}^{\text{likelihood}} \overbrace{p(\beta)}^{\text{prior}}}{\int p(y|\beta)p(\beta)d\beta} .$$

In this case, the estimate switches to from the maximum likelihood estimate to a summary of the location of the posterior distribution, such as the posterior median. The current literature on dealing with separation suggests researcher take an automatic approach by using a default prior distribution, such as Jeffreys’ invariant prior distribution (Jeffreys 1946 and Zorn 2005) or a heavy-tailed Cauchy(0, 2.5) prior distribution (Gelman et al. 2008).

Jeffreys’ Invariant Prior

Zorn (2005) suggests that political scientists deal with separation by maximizing a penalized

likelihood rather than the likelihood (see Heinze and Schemper 2002 as well). Zorn suggests replacing the usual likelihood function $L(\beta|y)$ with a “penalized” likelihood function $L^*(\beta|y)$, so that $L^*(\beta|y) = L(\beta|y)|I(\beta)|^{\frac{1}{2}}$. It turns out that the penalty $|I(\beta)|^{\frac{1}{2}}$ is equivalent to Jeffreys’ (1946) prior for the logistic regression model (Firth 1993 and Poirier 1994). Jeffreys’ prior can be obtained by applying Jeffreys’ Rule (Jeffreys 1946 and Box and Taio 2011, pp. 41-60), which requires setting the prior $p(\beta)$ to be proportional to the square root of the determinant of the information matrix, so that $p(\beta) = |I(\beta)|^{\frac{1}{2}}$. Then, of course, applying Bayes’ Rule yields the posterior distribution $p(\beta|y) \propto L(\beta|y)|I(\beta)|^{\frac{1}{2}}$, so that Firth’s penalized likelihood is equivalent to a Bayesian approach with Jeffreys’ prior. The researcher can then sample from this posterior distribution using MCMC to obtain the features of interest, such as the mean and standard deviation.

However, Firth (1993) did not propose this prior to solve the separation problem. Instead, his purpose was to reduce the well-known small sample bias in logistic regression models. And while it is true that Firth’s correction does provide finite estimates under separation, it remains an open question as to whether this automatic prior, designed for other purposes, provides a reasonable estimate of the uncertainty of the estimates for particular research problems.

The Cauchy(0, 2.5) Prior

Indeed, Gelman et al. (2008) note that Firth’s application of Jeffreys’ prior is not easily interpretable as an actual application of prior information because the prior $p(\beta) = |I(\beta)|^{\frac{1}{2}}$ lacks an interpretable scale and depends on the data in complex ways. Instead, they suggest standardizing continuous inputs to have mean zero and standard deviation one-half and simply centering binary inputs (Gelman 2008). Then, they suggest placing a weakly informative Cauchy(0, 2.5) prior on the coefficients for these rescaled variables that, like Jeffreys’ prior, bounds the estimates away from positive and negative infinity but can also be interpreted as actual prior

information.⁷ Gelman et al. (2008, p. 1363) write:

Our key idea is that actual effects tend to fall within a limited range. For logistic regression, a change of 5 moves a probability from 0.01 to 0.5, or from 0.5 to 0.99. We rarely encounter situations where a shift in input x corresponds to the probability of outcome y changing from 0.01 to 0.99, hence, we are willing to assign a prior distribution that assigns low probabilities to changes of 10 on the logistic scale.

As before, the posterior distribution is not easily available analytically, but one can easily use MCMC to simulate from the posterior distribution. Once a researcher has the MCMC simulations, she can obtain the point estimates and credible intervals for the parameters by summarizing the simulations.

Gelman et al. (2008) design their prior distribution to be reflective of prior informative for a range of situations. In many cases, their weakly informative prior supplies too little prior information. In a few cases, it might supply too much. In any either case, it remains an open question as to whether this general prior provides a reasonable estimate of the uncertainty of the estimates for *particular* research problems.

The Importance of the Prior

While default priors, such as Zorn's suggested Jeffreys' prior or Gelman et al.'s suggested Cauchy(0, 2.5) prior are often useful as starting points, choosing an informative prior distribution is crucial for dealing with separation in a substantively meaningful manner. Further, whether a particular prior is reasonable depends on the particular application.

In most data analyses, the data swamp the contribution of the prior, so that the choice of

⁷Gelman et al. (2008) use a Cauchy(0, 2.5) prior for the coefficients but a Cauchy(0, 10) prior for the *intercept*. This allows the intercept to take on a *much* larger range of values (e.g., from 10^{-9} to $1 - 10^{-9}$)

prior has little effect on the posterior. However, in the case of separation, the prior essentially determines the shape of the posterior in the direction of the separation. The likelihood has an “S”-shape that approaches a limit as the parameter coefficient for the separating variable s_i approaches infinity. Thus, for large values of the coefficient, the likelihood is essentially flat, which allows the prior distribution to drive the inferences. Thus, the prior distribution is not an arbitrary choice made for computational convenience—but an important choice that affects the inferences. We can see the importance in both theory and practice.

The Impact of the Prior in Theory

Although it is intuitive that the prior drives the inferences in the direction of the separation, it is also easy to characterize the impact of the prior on a monotonic increasing likelihood in a general way. Suppose quasicomplete separation, such that whenever an explanatory variable $s_i = 1$, a binary outcome $y_i = 1$, but when $s_i = 0$, y_i might equal zero or one. Suppose further that the analyst wishes to obtain plausible estimates of coefficients for the model

$$Pr(y_i = 1) = \text{logit}^{-1}(\beta_{cons} + \beta_s s_i + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_k x_{ik}).$$

It is easy to find plausible estimates for the coefficients of x_1, x_2, \dots, x_k using maximum likelihood, but finding a plausible estimate of β_s proves more difficult because maximum likelihood suggests an estimate of $+\infty$. In order to obtain a plausible estimate of β_s , the researcher must introduce prior information into the model. My purpose here is to characterize how this prior information impacts the posterior distribution.

In the general situation, the analyst is interested in computing and characterizing the posterior distribution of β_s given the data. Using Bayes’ Rule, the posterior distribution of $\beta = \langle \beta_{cons}, \beta_s, \beta_1, \beta_2, \dots, \beta_k \rangle$ depends on the likelihood and the prior, so that $p(\beta|y) \propto p(y|\beta)p(\beta)$. In particular, the analyst might have in mind a family of priors centered at and monotonically

decreasing away from zero with varying scale σ , so that $p(\beta_s) = p(\beta_s|\sigma)$, though the results below simply depend on having any proper prior distribution. The informativeness of the prior distribution depends on σ , which is chosen by the researcher and “flattens” the prior $p(\beta_s) = p(\beta_s|\sigma)$, such that as σ increases, the rate at which the prior descends to zero decreases.

Theorem 1. *For a monotonic likelihood $p(y|\beta)$ increasing [decreasing] in β_s , proper prior distribution $p(\beta|\sigma)$, and large positive [negative] β_s , the posterior distribution of β_s is proportional to the prior distribution for β_s , so that $p(\beta_s|y) \propto p(\beta_s|\sigma)$. More formally, $\lim_{\substack{\beta_s \rightarrow \infty \\ [-\infty]}} \frac{p(\beta_s|y)}{p(\beta_s|\sigma)} = k$, for positive constant k .*

Proof and details: See the Technical Appendix.

Theorem 1 simply implies that for large values of β_s the posterior distribution depends almost entirely on the researcher’s *choice* of prior distribution. Thus, the choice of prior matters. While the choice of prior might not affect the conclusion about the *direction* of the effect (i.e., one-sided credible interval), it has a large impact on the conclusion about the *magnitude* of the effect (i.e., two-sided credible interval).⁸ Credible intervals are crucial when discussing effect magnitudes (see King, Tomz, and Wittenberg 2000; Rainey 2014; Gross 2014; and McCaskey and Rainey 2014), and the choice of prior essentially drives the width of the credible interval.

The Impact of the Prior in Practice

To illustrate the impact of the prior on inferences when facing separation, I replicate results from Barrilleaux and Rainey (2014), who are interested in the effect of partisanship on governors’ decisions to oppose the Medicaid expansion in their states under the Patient Protection and Affordable Care Act (ACA).⁹ As the authors note, no Democratic governors opposed the

⁸Indeed, if one is only interested in the *direction* of an effect (i.e., “sign and significance”), then the procedure suggested by Geyer (2009) is a reasonable alternative to regularization with a prior and the choice of prior is not likely to make much difference.

⁹Barrilleaux and Rainey (2014) use a logistic regression modeling the probability that a governor opposes the expansion using the following explanatory variables: the partisanship of the governor, the percent of the state’s

expansion, which leads to a problem of separation. To see whether the choice of prior matters, I use MCMC to simulate from the posterior using Zorn's (2005) and Gelman et al.'s (2008) suggested *default* prior distributions.

Figure 1 shows the posterior medians and 90% credible interval for the two default priors.¹⁰ While the choice of prior does not affect the conclusion about the *direction* of the effect, it has a large impact on the conclusion about the *magnitude* of the effect. This can be especially important when researchers are making claims about the substantive importance of their estimated effects (see King, Tomz, and Wittenberg 2000; Rainey 2014; Gross 2014; and McCaskey and Rainey 2014). For example, the Cauchy(0, 2.5) prior leads to a posterior median that is over 40% larger than the posterior median from Jeffreys' prior (4.9 compared to 3.5). The posterior mean is more than 80% larger using the Cauchy(0, 2.5) prior (7.1 compared to 3.9). Further, the 90% credible interval is more than twice as wide for the Cauchy(0, 2.5) prior. The choice between two *default* priors leads to a large change in inferences.

residents who are favorable toward the ACA, whether Republicans control the state legislature, the percent of the state that is uninsured, a measure of the fiscal health of the state, the Medicaid multiplier for the state, the percent of the state that is nonwhite, and the percent of the state that resides in a metropolitan area. See their paper for more details.

¹⁰The credible intervals I use throughout this paper are 90% HDP credible intervals. Obviously, we could define many intervals that have 90% chance of containing the true parameter. However, the HPD interval is theoretically appealing because it is the *shortest* of these intervals. See Gill (2008, esp. pp. 48-51) and Casella and Berger (2002, esp. p. 448). The alternative, equal-tailed intervals, tends to exacerbate the differences between the priors.

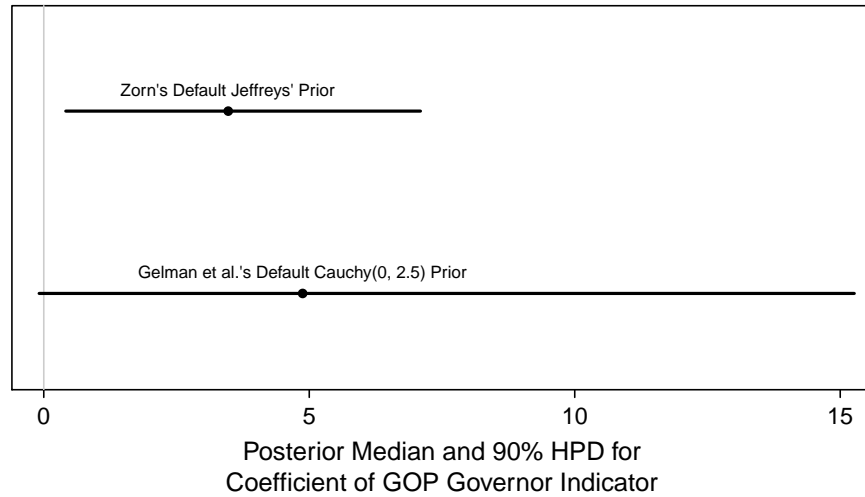


FIGURE 1: This figure provides the posterior medians and 90% credible intervals for the coefficient of the indicator for GOP governors in the model offered by Barrilleaux and Rainey (2014). Notice that Jeffreys' prior, suggested by Zorn (2005), is the more informative of these priors, suggesting that a coefficient larger than about seven is quite unlikely. On the other hand, the credible interval using the Cauchy(0, 2.5) prior, as suggested by Gelman et al. (2008), is about *twice* as wide as the credible interval from Jeffreys' prior, suggesting that effects as large as about 15 are plausible. Further, the posterior median from the Cauchy(0, 2.5) prior is about 40% larger than the posterior median from Jeffreys' prior.

Figure 2 shows the posterior distribution for the coefficient for the indicator of Republican governors. Notice that these two *default* priors lead to different posterior distributions. Notice, in particular, that the choice of the prior has a large impact on the right-hand side of the posterior, as suggested by Theorem 1. The more informative Jeffreys' prior leads to a more peaked posterior distribution that nearly rules out coefficients larger than about seven. The less informative Cauchy(0, 2.5) prior leads to the conclusion that much larger coefficients, such as 15, are plausible. These differences are not trivial—there are large differences in the posterior distributions, and these differences can affect the conclusions that the researchers draw about the likely magnitude of the effect.¹¹

¹¹But what about when the separating variable is not of direct interest? For example, what if the researcher simply wants to include a separating variable as a control? In this case, the choice of prior still changes the estimates, but not nearly as much. The substantive conclusions are not likely to depend on the choice of prior. It is *crucial* Zorn (2005) to include the variable in the model, but the prior distribution is not nearly as important as when the separating variable is of direct substantive interest.

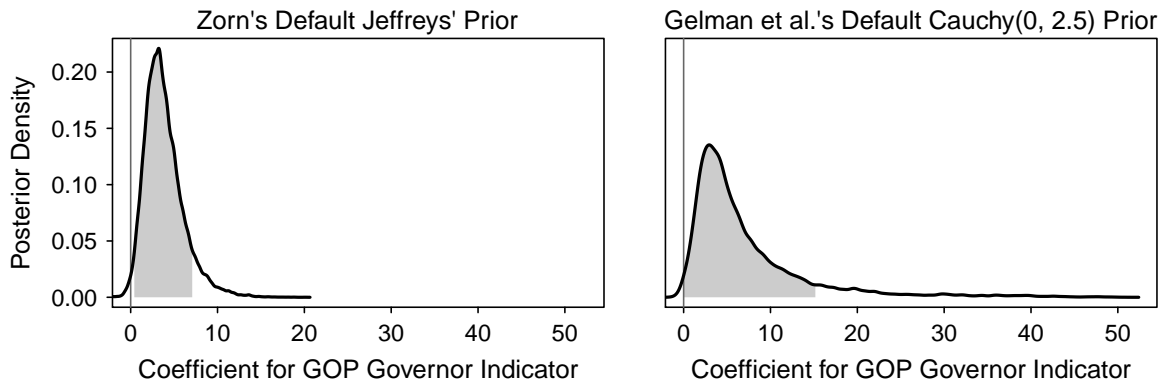


FIGURE 2: This figure shows the posterior distribution for the coefficient of the indicator for GOP governors in the model offered by Barrilleaux and Rainey (2014) for different default prior distributions. The grey shading indicates the 90% credible interval. Notice that the location and the spread of the posterior depend on the prior chosen, especially the right-hand side of the distribution, as suggested by Theorem 1.

Choosing an Informative Prior

While it is often sufficient to rely on default priors, this is not the case if one is interested in obtaining reasonable measures of uncertainty under separation. Indeed, in the replication of Barrilleaux and Rainey (2014) above, I show that the overall posterior distribution, the width of the 90% credible interval, and the posterior median largely depend on the prior one chooses. This implies that researchers relying on default priors alone risk under- or over-representing their confidence in the magnitude of the effect.

Under separation, the data simply offer no information about the upper-bound of the magnitude of the coefficient for the separating variable. Any reasonable regularization, then, must come in the form of an informative prior. Leamer (1983, p. 38) writes: “As I see it, the fundamental problem facing econometrics is how adequately to control the whimsical character of inference, how sensibly to base inferences on opinions when facts are unavailable.” However, a researcher’s prior need not be an spur-of-the-moment feeling, but contain information. This

information can come from several sources, including quantitative studies of similar topics, detailed analyses of particular cases, and theoretical arguments.

When facing separation, I suggest researchers use a prior distribution that satisfies three properties:

1. *Shrinks the estimates toward zero.* While the ultimate goal is to choose a prior distribution based on actual prior information, the prior distribution should also be appropriately conservative. As mentioned before, the prior distribution largely drives the inferences in the direction of the separation. In this case, a non-central prior distribution in the direction of the separation has an especially large impact on the inferences. For this reason, I focus on prior distributions centered at zero to conservatively shrink coefficients toward zero (Gelman and Jakulin 2007). The only choice the researcher needs to make is the amount of shrinkage appropriate for a given substantive problem.
2. *Allows plausible effects.* The prior distribution should assign substantial prior probabilities to estimates that are *a priori* plausible according to the researcher's own prior information.
3. *Rules out implausibly large effects.* The prior distribution should assign little prior probability to estimates that are *a priori* implausible according to the researcher's own prior information.¹²

Different researchers will inevitably have different prior beliefs. For example there is substantial disagreement among international relations theorists about the likely effects of nuclear weapons on conflict. Some optimists believe that nuclear weapons make peace much more likely. Mearsheimer (1993, p. 57) argues that “nuclear weapons are a powerful force for peace” and observes:

¹²Because all the priors I consider have support over the real line, they do not technically “rule out” anything. However, these priors rule out large effects *in practice*, by placing very little prior probability on large effects. Enough data will overwhelm these priors, but for typical social data sets, these priors can essentially rule out large effects.

In the pre-nuclear world of industrialized great powers, there were two world wars between 1900 and 1945 in which some 50 million Europeans died. In the nuclear age, the story is very different. Only some 15,000 Europeans were killed in minor wars between 1945 and 1990, and there was a stable peace between the superpowers that became increasingly robust over time. A principal cause of this “long peace” was nuclear weapons.

Bueno de Mesquita and Riker (1982, p. 283) even theorize that the probability of conflict “decreases to zero when all nations are nuclearly armed.” On the other hand, some pessimists (e.g., Sagan 1994) believe that nuclear weapons do not deter conflict, only make it more catastrophic. Mueller (1988, pp. 68-69) writes:

Nuclear weapons may well have enhanced this stability—they are certainly dramatic reminders of how horrible a big war could be. But it seems highly unlikely that, in their absence, the leaders of the major powers would be so unimaginative as to need such reminding. Wars are not begun out of casual caprice or idle fancy, but because one country or another decides that it can profit from (not simply win) the war—the combination of risk, gain, and cost appears to be preferable to peace. Even allow considerably for stupidity, ineptness, miscalculation, and self-deception in these considerations, it does not appear that a large war, nuclear or otherwise, has been remotely in the interest of essentially-contented, risk-averse, escalation-anticipating powers that have dominated world affairs since 1945.

The optimists and the pessimists have different prior beliefs about the likely effects nuclear weapons. These different beliefs must lead to different interpretations of the evidence because the prior distribution has such a strong impact on the posterior distribution in the direction of the separation. However, researchers must clearly communicate the dependence of the inferences on the choice of prior by transparently developing an informative prior distribution

and providing the inferences for alternative prior beliefs.

However, choosing a prior distribution is quite difficult, especially for multidimensional problems. Gill and Walker (2005) provide an overview of methods of choosing a prior appropriate to social science research. However, the most sensible approach for choosing a prior distribution depends on the nature of the statistical model and the prior information.

In general, the researcher might assess the reasonableness of the prior distribution by examining the prior predictive distribution and asking herself whether the prior and model produce a distribution for the data that matches her prior information. Under the Bayesian framework, the researcher has a fully specified model $p(y_{new}|\theta)p(\theta)$ and can thus simulate hypothetical data y_{new} from the model prior to observing the data. The distribution of the unobserved outcome y_{new} is given by $p(y_{new}) = \int p(y_{new}|\theta)p(\theta)d\theta$ (Box 1980). In practice, this process involves Clarify-like simulation (King, Tomz, and Wittenberg 2000), but rather than using the asymptotic posterior (e.g., $\beta_{sim} \sim N[\hat{\beta}^{mle}, I(\hat{\beta}^{mle})^{-1}]$), researchers simulate the model parameters from the prior distribution. Just as a researcher can use simulation to interpret the coefficient estimates of nonlinear models, she can use simulation to interpret the prior distribution.

Definition 1 (Prior Predictive Distribution). *The prior predictive distribution, denoted as $p(y_{new})$, is the prior distribution of hypothetical data, so that $p(y_{new}) = \int_{-\infty}^{\infty} p(y_{new}|\beta)p(\beta)d(\beta)$.*

However, it is difficult to work with more than one dimension of the prior distribution. Specifying the full prior distribution requires simultaneously choosing prior distributions for the $k + 1$ explanatory variables, as well as the relationships among these variables (e.g., family, location, scale, and correlations of all the parameters). This process is intractably tedious, as the researcher must evaluate the prior for each combination of each parameter set at a range of values. Even if the researcher considers only independent normal priors and ten values for each parameter, then the researcher must examine 10^{k+1} prior predictive distributions. If the researcher has eight control variables, so that $k = 8$, (e.g., Barrilleaux and Rainey 2014), then

the researcher must evaluate one *billion* prior predictive distributions.

But only specific regions of the $k + 1$ dimensional prior distribution are practically important when addressing separation. this allows the researcher to dramatically simplify the choice of prior. In particular, the researcher can simplify the focusing in two specific ways.

1. *Focus only on the separated coefficient.* Since the data swamp the prior for all the model coefficients except β_s , the only relevant “slices” of the prior distribution are those in which all other coefficients are near their maximum likelihood estimates.
2. *Focus in the direction of the separation.* The likelihood also swamps the prior in the direction opposite the separation. Unless the researcher has an extremely small data set (i.e., smaller than Barrilleaux and Rainey (2014), who have $N = 50$), the the likelihood essentially rules out values less [greater] than zero when the direction of separation is positive [negative].

I refer to this simplified focus as the *partial* prior predictive distribution.

Definition 2 (Partial Prior Predictive Distribution). *The partial prior predictive distribution, denoted as $p^*(y_{new})$, is the prior distribution of y_{new} given that the separated coefficient lies in the direction of the separation and all other coefficients equal their maximum likelihood estimates, so that $p^*(y_{new}) = \int_0^\infty p(y_{new}|\beta_s, \hat{\beta}_{-s}^{mle})p(\beta_s|\beta_s \geq 0)d(\beta_s)$ when $\hat{\beta}_s^{mle} = +\infty$ and $p^*(y_{new}) = \int_{-\infty}^0 p(y_{new}|\beta_s, \hat{\beta}_{-s}^{mle})p(\beta_s|\beta_s \leq 0)d(\beta_s)$ when $\hat{\beta}_s^{mle} = -\infty$.*

For example, in Barrilleaux and Rainey (2014), we do not need to use prior information to obtain reasonable estimates for our measures of need and public opinion. Further, because no Democratic governors opposed the Medicaid expansion, then we do not need the prior to rule out large *positive* effects for Democratic partisanship. In both cases, the likelihood is sufficiently informative.

However, we do need to use the prior to rule out large *negative* effects for Democratic partisanship, because the likelihood cannot effectively rule out implausibly large negative

effects. Indeed, the likelihood is monotonically decreasing in the coefficient for the indicator of Democratic governors. That is, the likelihood increases as the coefficient for Democratic partisanship becomes more negative. The larger the negative effect, the more likely separation would occur. The usual maximum likelihood, therefore, provides implausibly large negative estimates and unreasonable standard errors. Theorem 1 provides a more formal treatment of this intuition, but prior information is essential to obtain reasonable estimates and measures of uncertainty.

Choosing a prior, though, requires thoughtful effort. As I show above, default priors can lead to much different conclusions, so it is essential to build actual prior information into the model. In order to choose a reasonable, informative prior distribution, researchers need to obtain the partial prior predictive distribution defined in Definition 2. The following steps describe how researchers can simulate from the partial prior predictive distribution and use the simulations to check the reasonableness of the choice.

1. Estimate the model coefficients using maximum likelihood, given the coefficient vector $\hat{\beta}^{mle}$. Include the separating variable s_i in the model. Of course, this leads to implausible estimates for β_s , but the purpose is to choose reasonable values at which to fix the *other* parameters in order to focus on a single slice of the full prior.
2. Choose a prior distribution $p(\beta_s)$ for the separating variable s that is centered at zero. The most common choice is the scaled t distribution, which has the normal and Cauchy families as special cases ($df = \infty$ and $df = 1$, respectively).
3. Choose a large number of simulations n_{sims} to perform (e.g., $n_{sims} \geq 10,000$) and for i in 1 to n_{sims} , do the following:
 - a. Simulate $\tilde{\beta}_s^{[i]} \sim p(\beta_s)$.
 - b. Replace $\hat{\beta}_s^{mle}$ in $\hat{\beta}^{mle}$ with $\tilde{\beta}_s^{[i]}$, yielding the vector $\tilde{\beta}^{[i]}$.

- c. Calculate and store the quantity of interest $\tilde{q}^{[i]} = q(\tilde{\beta}^{[i]})$. This quantity of interest might be a first-difference or risk-ratio, for example.
4. Keep only those simulations in the direction of the separation (e.g., $\tilde{\beta}_s^{[i]} \geq 0$ when $\hat{\beta}_s^{mle} = +\infty$ and $\tilde{\beta}_s^{[i]} \leq 0$ when $\hat{\beta}_s^{mle} = -\infty$).
5. Summarize the simulations \tilde{q} using quantiles, histograms, or density plots. If the prior is inadequate, then update the prior distribution $p(\beta_s)$.

Given that the inference can be highly dependent on the choice of prior, I recommend that the researcher choose at least three prior distributions: (1) an *informative* prior distribution that represents her actual information, (2) a highly *skeptical* prior distribution that suggests the effect is likely small, and (3) a highly *enthusiastic* prior that represents the suggests the effect might be very large. Combined with Zorn's (2005) and Gelman et al.'s (2008) suggested defaults, these provide a range of prior distributions that researchers can use to evaluate their inferences.

Estimating the Full Model

Once the researcher obtains a reasonable prior distribution as well as several to use for robustness checks, she must use MCMC to obtain simulations from the posterior. Zorn (2005) and Gelman et al. (2008) suggest variations on maximum likelihood to quickly obtain estimates and confidence intervals. However, the normal approximation typically used to simulate the parameters and calculate quantities of interest (King, Tomz, and Wittenberg 2000) is particularly inaccurate under separation. As an alternative, I recommend the researcher use MCMC to simulate directly from the posterior distribution. The researcher can then use these simulations to calculate point estimates and confidence intervals for any desired quantity of interest. For the informative $p_{\text{inf}}(\beta_s)$, skeptical $p_{\text{skep}}(\beta_s)$, and enthusiastic $p_{\text{opt}}(\beta_s)$ priors, I suggest the model:

$$Pr(y_i = 1) = \text{logit}^{-1}(\beta_{cons} + \beta_s s_i + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_k x_{ik})$$

$$\beta_s \sim p_k(\beta_s), \text{ for } k \in \{\text{inf, skip, enth}\},$$

with improper, constant priors on the other model coefficients.

Application: Nuclear Proliferation and War

A recent debate emerged in the conflict literature between Rauchhaus (2009) and Bell and Miller (2015) that revolves around the issue of separation. Rauchhaus (2009, p. 262) hypothesizes that “[t]he probability of major war between two states will decrease if both states possess nuclear weapons.” Summarizing his empirical results, Rauchhaus writes:

The hypotheses on nuclear symmetry find strong empirical support. The probability of a major war between two states is found to decrease when both states possess nuclear weapons (p. 269).

Despite using the same data, Bell and Miller (2015, p. 9) claim that “nonnuclear dyads are in fact no more likely to fight wars than nonnuclear dyads.” Their disagreement hinges, in part, on whether and how to handle separation, because no nuclear dyad in Rauchhaus data engages in war.¹³ Rauchhaus (2009) ignores the separation and estimates that nonnuclear dyads are about 2.7 million times more likely to go to war than symmetric nuclear dyads. Bell and Miller (2015), on the other hand, use Jeffreys’ (1946) invariant prior, as suggested by Zorn (2005), and estimate that nonnuclear dyads are only about 1.6 times more likely to engage in war. Because these authors use very different prior distributions, they reach very different

¹³Bell and Miller (2015) also disagree with Rauchhaus’s (2009) coding of the 1999 conflict in Kargil between India and Pakistan, which both possessed nuclear weapons. This conflict is excluded from Rauchhaus’s data set, but Bell and Miller argue that it should be included as a war between two nuclear-armed states. However, that portion of their argument is less relevant to my purpose. Instead, my goal is to illustrate how one might choose a reasonable prior distribution and highlight the importance of the choice of prior.

conclusions.¹⁴ This raises important questions. First, would a reasonable, informative prior distribution support Rauchhaus' position of a meaningful effect or Bell and Miller's position of essentially no effect? Second, how robust is the conclusion to a range of more and less informative prior distributions?

Prior

The first step in dealing with the separation in a principled manner is to choose a prior distribution that represents actual prior information. To choose a reasonable prior, I follow the process above to generate a partial prior predictive distribution for the risk-ratio that Bell and Miller (2015) emphasize. I experimented with a range of prior distributions, from a variety of families, but settled on the normal family. After some experimentation, I selected a normal distribution with mean zero and standard deviation 4.5 to serve as an informative prior and represent my own prior beliefs. I chose this prior distribution because it essentially rules out risk-ratios larger than 1,000—effects that I find implausibly large—and treats risk-ratios smaller than 1,000 as plausible. Figure 3 and Table 1 summarize the partial prior predictive distributions for the this normal distribution with standard deviation 4.5.

To evaluate the robustness of any statistical claims to the choice of prior, I also selected a highly skeptical and highly enthusiastic prior. I chose a normal distribution with mean zero and standard deviation two to serve as a skeptical prior that represents the belief that any pacifying effect of nuclear weapons is small (e.g., Mueller 1988). This skeptical prior distribution essentially rules out risk-ratios larger than 25 as implausibly large. Finally, I selected a normal distribution with mean zero and standard deviation eight to serve as an enthusiastic prior that represents the belief that the pacifying effects of nuclear weapons might be quite large (e.g., Mearsheimer 1993). This enthusiastic prior, on the other hand, treats risk-ratios as large as

¹⁴Rauchhaus (2009) does not use a formal prior distribution, but uses generalized estimating equations, which we might interpret as having an improper, uniform prior on the logistic regression coefficients from minus infinity to plus infinity. The estimate is finite only due to a stopping rule in the iterative optimization algorithm.

500,000 as plausible. Figure 3 shows the partial prior predictive distributions for the informative, skeptical, and enthusiastic prior distributions. For convenience, Table 1 provides the deciles of the PPPDs shown in Figure 3.




FIGURE 3: This figure shows the partial prior predictive distribution for the risk-ratio of war in nonnuclear to nuclear dyads. The risk-ratio is tells us how many times more likely war is in non-nuclear dyads compared to nuclear dyads. Notice that the informative prior treats effects smaller than 100 as plausible, but essentially rules out effects larger than 10,000. The skeptical prior essentially rules out effects larger than 100, while the enthusiastic prior treats effects between 1 and 100,000 as essentially equally likely.

	10%	20%	30%	40%	50%	60%	70%	80%	90%
Informative Normal(0, 4.5) Prior	1.7	3.1	5.6	10.3	19.5	44.1	98.8	296.9	1,577.7
Skeptical Normal(0, 2) Prior	1.3	1.7	2.2	2.9	3.9	5.4	8.3	13.2	26.7
Enthusiastic Normal(0, 8) Prior	2.9	8.1	24.7	76.5	256	987.6	5,300.9	42,466.7	643,954.6

TABLE 1: This table provides the deciles prior predictive distribution for the risk-ratio of war in nonnuclear and nuclear dyads. The risk-ratio is tells us how many times more likely war is in non-nuclear dyads compared to nuclear dyads. Notice that that the, informative prior suggests a median risk-ratio of about 20, which is a large, but plausible effect. The skeptical prior suggests a median ratio of about 4 and the enthusiastic prior suggests a median ratio of over 200.

Notice that the skeptical prior suggests that risk-ratios above and below 4 are equally likely (i.e., 50th percentile of the PPPD is 3.9), while the enthusiastic prior suggests that effects above and below 220 are equally likely. The informative prior, on the other hand, suggests (more reasonably, in my view) that the effect is equally likely to fall above and below 20. These three prior distributions, along with the defaults suggested by Zorn (2005) and Gelman et al. (2008), provide a range of distributions to represent a range of prior beliefs.

Posterior

Figure 4 shows the posterior distributions for the coefficient of the indicator of nuclear dyads from the informative, skeptical, enthusiastic, and two default prior distributions. The areas shaded grey indicate the 90% credible intervals. Notice that the location (e.g., peak or mode), shape, scale, and credible interval depend on the choice of prior. While the magnitude of this coefficient is not easily interpretable, notice that Gelman et al.'s (2008) suggested default prior is somewhat similar to the informative prior, but Zorn's (2005) suggested default is quite similar to the *skeptical* prior. Thus, these distributions illustrate that the prior is important when dealing with separation. Indeed, it is a critical step in obtaining reasonable inferences.

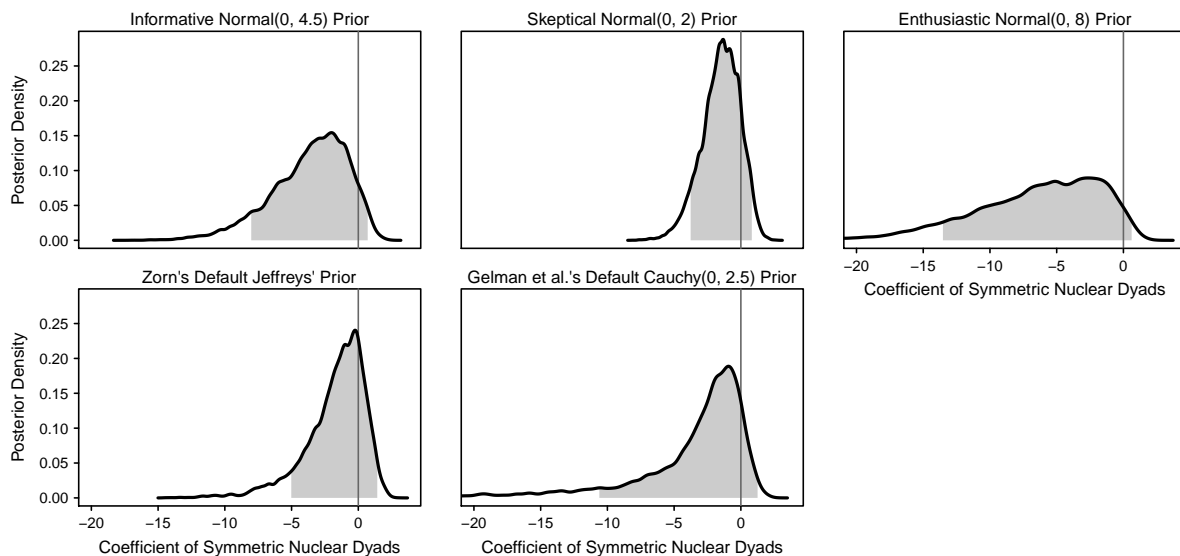


FIGURE 4: This figure shows the posterior distribution for the logit coefficient using each of the five prior distributions. The grey shading indicates the 90% credible interval. Notice that the choice of prior has a large effect on the inferences. For example, the enthusiastic prior suggests the ratio might be as large as -13, while the skeptical prior suggests the ratio might be as large as -4. Importantly, notice the similarity between the posterior from Zorn's (2005) suggested default and the skeptical prior, in terms of their peak (i.e., mode), shape, and highest posterior density.

I now turn to the posterior distribution of the risk-ratios—the key quantity of interest in the debate between Bell and Miller (2015) and Rauchhaus (2009). Figure 5 presents the posterior

medians and the 90% credible intervals for each prior. An initial glance at the figure shows substantial variation in the point estimates and the width of the intervals. However, these risk-ratios and credible intervals are plotted on the log scale (otherwise the wider intervals dominate the plot), so the figure tends to *understate* the variation in the inferences across priors. Notice that the informative prior suggests the true risk-ratio has about a 90% chance of falling between about 0.1 and about 2000, with a posterior median of about 25.¹⁵ The skeptical prior, on the other hand, suggests the risk-ratio has about a 90% chance of falling between 0.1 and 30, with a posterior median of about 4. The enthusiastic prior suggests the risk-ratio has about a 90% chance of falling between 0.1 and 500,000. The inferences from these priors are *very* different.

Further, the 90% credible interval using Zorn's (2005) default is *much* narrower than the informative prior, and the posterior median of Zorn's suggested default prior is even smaller than the *skeptical* prior. For this particular application, Gelman's suggested default more closely matches the informative prior, though notice that the point estimate from Gelman's prior is 62% smaller than that from the informative prior and the upper-bound of the credible interval is 13 times larger.

¹⁵It is worth pointing out that these wide confidence interval suggests that, although the credible intervals overlap zero, these data do not warrant a conclusion of "no effect" (see Rainey 2014).

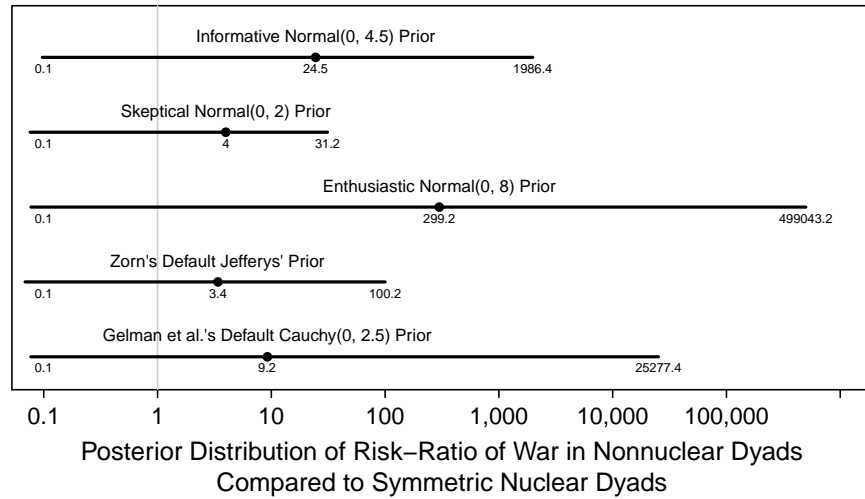


FIGURE 5: This figure shows the posterior median and 90% credible intervals for the risk-ratio using each of the five prior distributions *on the log scale*. Notice that the choice of prior has a huge effect on the inferences about the risk-ratio. For example, the skeptical prior suggests the ratio might be as large as 31, while the enthusiastic prior suggests the ratio might be as large as 500,000. Also, notice that the posterior median from Zorn's proposed default prior is *smaller* than the posterior median from the skeptical prior.

Finally, I use the posterior distributions from each prior to calculate the probability that the presence of nuclear weapons make war less likely (i.e., the probability that the risk-ratios shown in Figure 5 are greater than one). Recall Rauchhaus' hypothesis that nuclear weapons decrease the chance of war. These probabilities can be thought of as the probability that Rauchhaus' hypothesis is correct. Following the standard of $p \leq 0.05$ as offering strong evidence against the null hypothesis, it is reasonable to take $Pr(RR > 1) \geq 0.95$ as strong evidence for the research hypothesis. Figure 6 shows the probability that the hypothesis is correct for each prior distribution. Notice that while only the enthusiastic prior falls above the 0.95 standard, the evidence for the claim is at least suggestive. Perhaps most importantly for my purposes, Zorn's suggested default leads to the *least* evidence in favor of Rauchhaus' hypothesis—even the skeptical prior provides more evidence for Rauchhaus' claim.¹⁶

¹⁶Bell and Miller (2015) also question Rauchhaus' (2009) coding of the 1999 conflict in Kargil between India and Pakistan. If this conflict is coded as a war, as Bell and Miller suggest, then separation no longer occurs. With




FIGURE 6: This figure shows the posterior probability of the hypothesis that nonnuclear dyads are *more* likely to engage in war than symmetric nuclear dyads for each of the five prior distribution. From a hypothesis testing perspective, the evidence for the hypothesis is borderline or suggestive for each prior. However, notice that the skeptical prior, perhaps held by a researcher who believes the pacifying effect of nuclear weapons is small or nil, yields *greater* evidence for the hypothesis than Jeffreys' invariant prior suggested as a default by Zorn (2005).

Conclusion

Separation is a relatively common situation in political science. It is also an unusual “problem” because the effects in the sample are “too big” for maximum likelihood. In this situation, dropping the separating variables (i.e., deliberate specification bias) or interpreting the implausible coefficients and standard errors are particularly unattractive options. But even the most principled solution to date, the incorporation of prior information via default priors (Zorn 2005 and Gelman et al. 2008), has shortcomings.

First, the normal approximations necessary to simulate from these default prior distributions perform poorly. While it is possible to use profile likelihood methods to obtain more accurate confidence intervals for the coefficients (Zorn 2005, Heinze and Schemper 2002, and McCullagh this different coding, the data *do* offer information about the likely upper-bound on the magnitude of coefficient for the indicator of nuclear dyads, and the different priors produce similar inferences.

and Nelder 1989), it is difficult to translate these intervals into confidence intervals for quantities of interest. I provide the computational tools to simulate directly from the posterior using both Zorn's (2005) and Gelman et al.'s (2008) suggested default priors.

Second, the applications examining the effect of nuclear weapons and the effect of governors' partisanship illustrate what Theorem 1 proves—under separation, the choice of prior affects substantive conclusions. Even the predominant default priors used to deal with separation can provide very different inferences. A carefully-chosen, informative prior is an essential step in the process of obtaining reasonable inferences under separation. But what does this mean for applied researchers? I suggest two implications:

1. When facing separation, the choice of prior matters. Researchers must carefully choose a prior that represents actual prior information. Otherwise, the point and interval estimates will be too small or too large.
2. In addition to carefully choosing an informative prior that represents her own beliefs, the researcher should show how the inferences change for a range of prior distributions. In the debate between Bell and Miller (2015) and Rauchhaus (2009), the choice of prior almost completely drives the inferences about the likely magnitude of the risk-ratio. Thus, to the extent that there is disagreement about the prior, there will be disagreement about the results. In particular, I suggest that researchers report the key quantities of interest for a skeptical prior, an enthusiastic prior, and the two default priors suggested by Zorn (2005) and Gelman et al. (2008).

When facing separation, researchers must *carefully* choose a prior distribution to nearly rule out implausibly large effects. This paper introduces the concept of a partial prior predictive distribution and the associated computational tools to help researchers choose a prior distribution that represents actual prior information for their particular research problem. By presenting results using several prior distributions, including an informative prior, researchers can increase

the transparency, credibility, and accuracy of their inferences when dealing with separation.

References

- Ahlquist, John S. 2010. "Building Strategic Capacity: The Political Underpinning of Coordinated Wage Bargaining." *American Political Science Review* 104(1):171–188.
- Albert, A., and J. A. Anderson. 1984. "On the Existence of Maximum Likelihood Estimates in Logistic Regression Models." *Biometrika* 71(1):1–10.
- Barrilleaux, Charles, and Carlisle Rainey. 2014. "The Politics of Need: Examining Governors' Decisions to Oppose the 'Obamacare' Medicaid Expansion." Forthcoming in *State Politics and Policy Quarterly*. Manuscript at <http://www.carlislerainey.com/files/need.pdf>.
- Bell, Mark S., and Nicholas L. Miller. 2015. "Questioning the Effect of Nuclear Weapons on Conflict." *Journal of Conflict Resolution* 59(1):74–92.
- Berry, Frances Stokes, and William D. Berry. 1990. "State Lottery Adoptions as Policy Innovations: An Event History Analysis." *American Political Science Review* 84(2):395–415.
- Box, George E. P. 1980. "Sampling and Bayes' Inference in Scientific Modelling and Robustness." *Journal of the Royal Statistical Society A* 143(4):383–430.
- Box, George E. P., and George C. Taio. 2011. *Bayesian Inference in Statistical Analysis*. New York: John Wiley and Sons.
- Brown, Robert L., and Jeffrey M. Kaplow. 2014. "Talking Peace: IAEA Technical Cooperation and Nuclear Proliferation." *Journal of Conflict Resolution* 58(3):402–428.
- Bueno de Mesquita, Bruce, and William Riker. 1982. "An Assessment of the Merits of Selec- An

- Assessment of the Merits of Selective Nuclear Proliferation.” *Journal of Conflict Resolution* 26(2):283–306.
- Casella, George, and Roger L. Berger. 2002. *Statistical Inference*. 2nd ed. Pacific Grove, California: Duxbury.
- Casellas, Jason P. 2009. “Coalitions in the House? The Election of Minorities to State Legislatures and Congress.” *Political Research Quarterly* 62(1):120–131.
- Cederman, Lars-Erik, Kristian Skrede Gleditsch, and Simon Hug. 2013. “Elections and Ethnic Civil War.” *Comparative Political Studies* 46(3):387–417.
- Cox, Gary W., Thad Kousser, and Matthew D. McCubbins. 2010. “Party Power or Preferences? Quasi-Experimental Evidence from American State Legislatures.” *Journal of Politics* 72(3):799–811.
- DeRouen, Jr., Karl R., and Jacob Bercovitch. 2008. “Enduring Internal Rivalries: A New Framework for the Study of Civil War.” *Journal of Peace Research* 45(1):55–74.
- Desposato, Scott, and Ethan Scheiner. 2008. “Governmental Centralization and Party Affiliation: Legislator Strategies in Brazil and Japan.” *American Political Science Review* 102(4):509–524.
- Fearon, James D. 1994. “Signaling versus the Balance of Power and Interests: An Empirical Test of a Crisis Bargaining Model.” *Journal of Conflict Resolution* 38(2):236–269.
- Firth, David. 1993. “Bias Reduction of Maximum Likelihood Estimates.” *Biometrika* 80(1):27–38.
- Fuhrmann, Matthew. 2012. “How “Atoms for Peace” Programs Cause Nuclear Insecurity.”
- Gelman, Andrew. 2008. “Scaling Regression Inputs by Dividing by Two Standard Deviations.” *Statistics in Medicine* 27(15):2865–2873.

- Gelman, Andrew, Aleks Jakulin, Maria Grazia Pittau, and Yu-Sung Su. 2008. "A Weakly Informative Prior Distribution for Logistic and Other Regression Models." *The Annals of Applied Statistics* 2(4):1360–1383.
- Gelman, Andrew, and Aleks Jakulin. 2007. "Bayes: Radical, Liberal, or Conservative?" *Statistica Sinica* 17(2):422–426.
- Geyer, Charles J. 2009. "Likelihood Inference in Exponential Families and Directions of Recession." *Electronic Journal of Statistics* 3:259–289.
- Gill, Jeff. 2008. *Bayesian Methods: A Social and Behavioral Science Approach*. 2nd ed. Boca Raton, FL: Chapman and Hall.
- Gill, Jeff, and Lee D. Walker. 2005. "Elicited Priors for Bayesian Model Specifications in Political Science Research." *Journal of Politics* 67(3):841–872.
- Gross, Justin H. 2014. "Testing What Matters (If You Must Test at All): A Context-Driven Approach to Substantive and Statistical Significance." Forthcoming in *American Journal of Political Science*.
- Heinze, Georg, and Michael Schemper. 2002. "A Solution to the Problem of Separation in Logistic Regression." *Statistics in Medicine* 21(16):2409–2419.
- Heller, William B., and Carol Mershon. 2008. "Dealing in Discipline: Party Switching and Legislative Voting in the Italian Chamber of Deputies, 1988-2000." *American Journal of Political Science* 52(4):910–924.
- Jeffreys, H. 1946. "An Invariant Form of the Prior Probability in Estimation Problems." *Proceedings of the Royal Society of London, Series A* 186(1007):453–461.

- King, Gary, Michael Tomz, and Jason Wittenberg. 2000. "Making the Most of Statistical Analyses: Improving Interpretation and Presentation." *American Journal of Political Science* 44(2):341–355.
- Leamer, Edward E. 1983. "Lets Take the Con Out of Econometrics." *American Economic Review* 73(1):31–43.
- Leeman, Lucas, and Isabella Mares. 2014. "The Adoption of Proportional Representation." *Journal of Politics* 76(2):461–478.
- Lesaffre, E., and A. Albert. 1989. "Partial Separation in Logistic Discrimination." *Journal of the Royal Statistical Society. Series B (Methodological)*. 51(1):109–116.
- Mares, Isabela. 2015. *From Open Secrets to Secret Voting: Democratic Electoral Reforms and Voter Autonomy*. Cambridge Studies in Comparative Politics Cambridge: Cambridge University Press.
- Martin, Lanny W., and Randolph T. Stevenson. 2001. "Government Formation in Parliamentary Democracies." *American Journal of Political Science* 45(1):33–50.
- McCaskey, Kelly, and Carlisle Rainey. 2014. "Meaningful Inferences: The Importance of Explicit Statistical Arguments for Substantive Significance." Working paper. Latest version at <https://github.com/carlislerainey/meaningful-inferences>.
- McCullagh, Peter, and John A. Nelder. 1989. *Generalized Linear Models*. Second ed. Boca Raton, FL: Chapman and Hall.
- Mearsheimer, John J. 1993. "The Case for a Ukrainian Nuclear Deterrent." *Foreign Affairs* 72(3):50–66.
- Minozzi, William, and Craig Volden. 2013. "Who Heeds the Coll of the Party in Congress?" *Journal of Politics* 75(3):787–802.

- Mueller, John. 1988. "The Essential Irrelevance of Nuclear Weapons: Stability in the Postwar World." *International Security* 13(2):55–79.
- Peterson, Timothy M., and A. Cooper Drury. 2011. "Sanctioning Violence: The Effect of Third-Party Economic Coercion on Militarized Conflict." *Journal of Conflict Resolution* 55(4):580–605.
- Poirier, Dale. 1994. "Jeffreys' Prior for Logit Models." *Journal of Econometrics* 63(2):327–339.
- Rainey, Carlisle. 2014. "Aruging for a Negligible Effect." Forthcoming in *American Journal of Political Science*.
- Rauchhaus, Robert. 2009. "Evaluating the Nuclear Peace Hypothesis: A Quantitative Approach." *Journal of Conflict Resolution* 53(2):258–278.
- Reiter, Dan. 2014. "Security Commitments and Nuclear Proliferation." *Foreign Policy Analysis* 10(1):61–80.
- Rocca, Michael S., Gabriel R. Sanchez, and Jason L. Morin. 2011. "The Institutional Mobility of Minority Members of Congress." *Political Research Quarterly* 64(4):897–909.
- Sagan, Scott D. 1994. "The Perils of Proliferation." *International Security* 18(4):66–107.
- Smith, Daniel A., and Dustin Fridkin. 2008. "Delegating Direct Democracy: Interparty Legislative Competition and the Adoption of the Initiative in the American States." *American Political Science Review* 102(333-350).
- Vining, Jr., Richard L., Teena Wilhelm, and Jack D. Collens. 2015. "A Market-Based Model of State Supreme Court News: Lessons from Captial Cases." *State Politics and Policy Quarterly* 15(1):3–23.

- Weisiger, Alex. 2014. "Victory Without Peace: Conquest, Insurgency, and War Termination." *Conflict Management and Peace Science* 31(4):357–382.
- Wolfinger, Raymond E., and Steven J. Rosenstone. 1980. *Who Votes?* Yale University Press.
- Zorn, Christopher. 2005. "A Solution to Separation in Binary Response Models." *Political Analysis* 13(2):157–170.

Technical Appendix

Proof of Theorem 1

Recall Theorem 1:

Theorem 1. For a monotonic likelihood $p(y|\beta)$ increasing [decreasing] in β_s , proper prior distribution $p(\beta|\sigma)$, and large positive [negative] β_s , the posterior distribution of β_s is proportional to the prior distribution for β_s , so that $p(\beta_s|y) \propto p(\beta_s|\sigma)$. More formally, $\lim_{\substack{\beta_s \rightarrow \infty \\ [-\infty]}} \frac{p(\beta_s|y)}{p(\beta_s|\sigma)} = k$, for positive constant k .

Proof. Due to separation, $p(y|\beta)$ is monotonic increasing in β_s to a limit $\underline{\mathcal{L}}$, so that $\lim_{\beta_s \rightarrow \infty} p(y|\beta_s) = \underline{\mathcal{L}}$. By Bayes' rule,

$$p(\beta|y) = \frac{p(y|\beta)p(\beta|\sigma)}{\int_{-\infty}^{\infty} p(y|\beta)p(\beta|\sigma)d\beta} = \frac{p(y|\beta)p(\beta|\sigma)}{\underbrace{p(y|\sigma)}_{\text{constant w.r.t. } \beta}}.$$

Integrating out the other parameters $\beta_{-s} = \langle \beta_{cons}, \beta_1, \beta_2, \dots, \beta_k \rangle$ to obtain the posterior distribution of β_s ,

$$p(\beta_s|y) = \frac{\int_{-\infty}^{\infty} p(y|\beta)p(\beta|\sigma)d\beta_{-s}}{p(y|\sigma)}, \quad (1)$$

and the prior distribution of β_s ,

$$p(\beta_s|\sigma) = \int_{-\infty}^{\infty} p(\beta|\sigma)d\beta_{-s}.$$

Notice that $p(\beta_s|y) \propto p(\beta_s|\sigma)$ iff $\frac{p(\beta_s|y)}{p(\beta_s|\sigma)} = k$, where the constant $k \neq 0$. Thus, Theorem 1

implies that

$$\lim_{\beta_s \rightarrow \infty} \frac{p(\beta_s|y)}{p(\beta_s|\sigma)} = k$$

Substituting in Equation 1,

$$\lim_{\beta_s \rightarrow \infty} \frac{\frac{\int_{-\infty}^{\infty} p(y|\beta)p(\beta|\sigma)d\beta_{-s}}{p(y|\sigma)}}{p(\beta_s|\sigma)} = k.$$

Multiplying both sides by $p(y|\sigma)$, which is constant with respect to β ,

$$\lim_{\beta_s \rightarrow \infty} \frac{\int_{-\infty}^{\infty} p(y|\beta)p(\beta|\sigma)d\beta_{-s}}{p(\beta_s|\sigma)} = kp(y|\sigma).$$

Setting $\int_{-\infty}^{\infty} p(y|\beta)p(\beta|\sigma)d\beta_{-s} = p(y|\beta_s)p(\beta_s|\sigma)$,

$$\lim_{\beta_s \rightarrow \infty} \frac{p(y|\beta_s)p(\beta_s|\sigma)}{p(\beta_s|\sigma)} = kp(y|\sigma).$$

Canceling $p(\beta_s|\sigma)$ in the numerator and denominator,

$$\lim_{\beta_s \rightarrow \infty} p(y|\beta_s) = kp(y|\sigma).$$

Recalling that $\lim_{\beta_s \rightarrow \infty} p(y|\beta) = \underline{\mathcal{L}}$ and substituting,

$$\underline{\mathcal{L}} = kp(y|\sigma),$$

which implies that $k = \frac{p(y|\sigma)}{\underline{\mathcal{L}}}$, which is a positive constant. □